

Contents

Enphase Monitor	2
Table of Contents	2
Quick Start	2
Features	3
Prerequisites	3
Installation	3
Step 1: Clone the Repository	3
Step 2: Install Dependencies	3
Step 3: Build the Application	3
Configuration	3
Initial Setup	3
Configuration Sections	4
Finding Your System IDs	5
OAuth Setup	5
Quick Setup	5
Detailed Guide	5
After OAuth Setup	6
API Configuration	6
What the Cloud API Provides	6
Rate Limits	6
Getting API Credentials	6
Usage	6
Run Once (Single Query)	6
Continuous Monitoring	7
Command-Line Options	7
Examples	7
Output Format	7
Metrics Explained	8
Combined Energy Report	8
Individual System Metrics	9
Color Customization	9
Troubleshooting	9
“API configuration required”	9
“API request failed with status 401”	9
“API request failed with status 404”	10
“rate limit exceeded (429)”	10
“API request failed with status 422”	10
No telemetry data returned	10
Caching and Rate Limits	10
Rate Limits	10
Caching Strategy	10
Cache File Format and Naming	11
Cache Management Commands	11
Best Practices	11
Documentation	12
For Getting Started	12

For Understanding the Codebase	12
Recommended Learning Path	12
Project Structure	12
Testing	14
Test Coverage by Package	14
Running Tests	14
Lint and CI	15
Test Mode (Cache Only)	15
Setting Up Test Data	16
Expected Values Format	16
Validation Tolerance	17
Performance	17
Code Quality	17
License	17

Enphase Monitor

A Go application for monitoring and aggregating data from multiple Enphase solar systems via the Enphase Enlighten Cloud API v4.

New to Go? This codebase follows Go best practices and includes comprehensive documentation. See **GO_BEST_PRACTICES.md** for a guide to Go patterns and idioms, and **GO_CONCEPTS.md** for a reference of Go concepts used in the code.

Table of Contents

- Quick Start
- Features
- Installation
- Configuration
- OAuth Setup
- Usage
- Output Format
- API Configuration
- Caching and Rate Limits
- Testing
- Troubleshooting
- Documentation
- Project Structure

Quick Start

New to this project? Start here:

1. **QUICKSTART.md** - Get up and running in 5 minutes
2. **OAuth_SETUP.md** - Complete OAuth authentication (required for first-time setup)
3. Return here for detailed usage and configuration

Features

- **Multi-System Monitoring:** Query and combine metrics from multiple independent Enphase systems
- **Comprehensive Metrics:** Track production, consumption, battery usage, grid import/export, and net energy flow
- **Flexible Querying:** Query historical dates or monitor real-time with auto-refresh
- **Clean Display:** Formatted terminal output with customizable colors
- **API Caching:** Automatic caching of API responses to reduce API calls and enable offline validation
- **Color Customization:** Customize terminal colors using hex codes or ANSI escape codes
- **Cloud API v4:** Uses Enphase Enlighten Cloud API v4 for reliable data access
- **OAuth Support:** Full OAuth 2.0 support for developer plan authentication
- **Test Mode:** Validation mode for testing against expected values without making API calls

Prerequisites

- Go 1.21 or higher
- **Enphase Developer Portal Account:** Register at <https://developer-v4.enphase.com/>
- **OAuth Credentials:** Get `api_key`, `client_id`, and `client_secret` from the Developer Portal
- **System IDs:** Your Enphase system IDs (find these from Enlighten URLs - see “Finding Your System IDs” section)

Installation

Step 1: Clone the Repository

```
git clone <repository-url>
cd enphase-monitor
```

Or download and extract the project files.

Step 2: Install Dependencies

```
go mod download
```

Step 3: Build the Application

```
go build -o enphase-monitor
```

Configuration

Initial Setup

1. Copy the example configuration file:

```
cp config.yaml.example config.yaml
```

2. Edit `config.yaml` with your credentials:

```
api:
  key: "YOUR_API_KEY"  # API Key from Enphase Developer Portal
```

```

client_id: "YOUR_CLIENT_ID" # OAuth Client ID
client_secret: "YOUR_CLIENT_SECRET" # OAuth Client Secret
authorization_url: "https://api.enphaseenergy.com/oauth/token"
redirect_uri: "http://localhost:8080/callback"
refresh_token: "YOUR_REFRESH_TOKEN" # From OAuth setup (see below)

systems:
  - name: "Left Subpanel"
    id: "YOUR_SYSTEM_ID" # Enlighten system ID

  - name: "Right Subpanel"
    id: "YOUR_OTHER_SYSTEM_ID" # Different system ID

refresh_interval: 3600 # Refresh interval in seconds (1 hour recommended)

# Optional: Timezone for reporting/display
# If not specified, uses OS system timezone (or US/Pacific if OS is UTC)
timezone: "US/Pacific" # Examples: "US/Pacific", "America/New_York", "Europe/London"

# Optional: Color customization
colors:
  production: "#f0b57c" # Solar Production
  discharge: "#7acf38" # Battery Discharge
  import: "#f63cb1" # Grid Import
  export: "#06b6de" # Grid Export
  net_import: "#f63cb1" # Net Energy Flow (Import)
  net_export: "#06b6de" # Net Energy Flow (Export)
  headers: "#f37320" # Report Headers
  charge: "#7acf38" # Battery Charge
  total_consumed: "#f37320" # Total Consumed
  secondary_text: "#808080" # Secondary Text
  primary_text: "#ffffff" # Primary Text
  error: "#ff0000" # Error Text

```

Implementation Details: See config.go for color conversion logic (hex → ANSI) and display.go for color usage in terminal output.

Configuration Sections

API Credentials You will need these from the Enphase Developer Portal: - `api.key`: Your API key - `api.client_id`: OAuth client ID - `api.client_secret`: OAuth client secret - `api.refresh_token`: Obtained from OAuth setup (see OAuth Setup below)

System Configuration Each system requires: - `name`: A friendly name for display - `id`: Your Enphase system ID (see Finding Your System IDs)

Optional Settings

- **refresh_interval**: How often to query the API in continuous mode (default: 3600 seconds = 1 hour)
 - **Important**: Only applies when running in continuous mode (without `--once` flag)
 - **Recommended**: Use 3600 seconds (1 hour) to respect API rate limits
 - **Rate Limit Consideration**: The API allows 10 requests per minute. With multiple systems, a low **refresh_interval** (e.g., 5 seconds) can quickly exceed this limit. For 2 systems, you would make 24 requests per minute (2 systems × 12 queries/minute), which exceeds the 10/minute limit.
 - **Best Practice**: Use 3600 seconds (1 hour) or higher to stay well within rate limits
- **timezone**: Timezone for reporting and display (optional)
 - **Default**: Uses your OS system timezone, or US/Pacific if OS timezone is UTC
 - **Format**: IANA timezone identifier (e.g., "US/Pacific", "America/New_York", "Europe/London")
 - **Examples**:


```
timezone: "US/Pacific"           # Pacific Time
timezone: "America/New_York"    # Eastern Time
timezone: "Europe/London"       # GMT/BST
```
 - **Usage**: All date/time display and API queries use this timezone. If not specified, the application uses your OS system timezone, falling back to US/Pacific if the system timezone is UTC.
- **colors**: Customize terminal output colors (see Color Customization)

Finding Your System IDs

To find your System IDs:

1. Log into <https://enlighten.enphaseenergy.com>
2. Select one of your systems
3. Look at the URL - it will be: `https://enlighten.enphaseenergy.com/systems/SYSTEM_ID/overview`
4. The number in the URL is your System ID
5. Repeat for each system you want to monitor

OAuth Setup

Required for First-Time Setup

This application uses OAuth 2.0 for authentication. You must complete a one-time OAuth setup to get your refresh token.

Quick Setup

```
./enphase-monitor --setup-oauth
```

Run the interactive setup wizard that will guide you through: 1. Generating an authorization URL 2. Authorizing the application in your browser 3. Exchanging the authorization code for tokens 4. Adding the refresh token to your config

Detailed Guide

For a comprehensive explanation of OAuth 2.0, what each component does, and how authentication works, see:

OAuth_SETUP.md - Complete OAuth guide with: - Explanation of OAuth 2.0 concepts - What the API server expects for authentication - How authorization works - Step-by-step setup instructions - Troubleshooting common issues

After OAuth Setup

Once you have your `refresh_token`, add it to your `config.yaml`:

```
api:
  refresh_token: YOUR_REFRESH_TOKEN  # From OAuth setup
```

The application will automatically use this token to get new access tokens when needed.

API Configuration

This application uses the **Enphase Enlighten Cloud API v4** exclusively to fetch energy data.

What the Cloud API Provides

- **Direct Daily Values:** The API returns pre-calculated daily totals - no need for state files or cumulative meter tracking
- **Historical Data:** Query any past date (not limited to today)
- **Real-time Updates:** Data is typically updated every 15 minutes
- **Reliable Access:** Works from anywhere with internet (no local network required)
- **Standardized Format:** Consistent JSON responses across all system types

Rate Limits

The free developer plan has the following limits: - **10 requests per minute** per API key - **1000 requests per month** total

The application automatically caches API responses to minimize API calls. The default refresh interval is 1 hour (3600 seconds) to stay well within these limits.

Getting API Credentials

1. **Register:** Create an account at <https://developer-v4.enphase.com/>
2. **Create Application:** Register a new application to get API credentials
3. **Get Credentials:** You will receive:
 - `api_key`: Your API key (used in all requests)
 - `client_id`: OAuth client ID
 - `client_secret`: OAuth client secret
4. **Complete OAuth Setup:** See OAuth Setup section above

Usage

Run Once (Single Query)

Query today's data and exit:

```
./enphase-monitor --once
```

Query a specific historical date:

```
./enphase-monitor --once --date 2026-01-15
```

Note: When querying a past date, the program automatically runs once (even without `--once` flag) since historical data doesn't change over time. You'll see a message like:

Note: Running once for historical date 2026-01-15 (data won't change)

Continuous Monitoring

Monitor with auto-refresh (uses `refresh_interval` from config):

```
./enphase-monitor
```

The application will query all systems at the configured `refresh_interval` (default: 3600 seconds = 1 hour) and display updated metrics.

Press `Ctrl+C` to stop.

Command-Line Options

- `--config <path>` - Path to configuration file (default: `config.yaml`)
- `--once` - Run once and exit instead of continuous monitoring
- `--date <YYYY-MM-DD>` - Query specific date instead of today (e.g., 2026-01-15)
- `--setup-oauth` - Run OAuth setup wizard (one-time for developer plan)
- `--test` - Test mode: use cache only, no live API calls, validate against expected values
- `--no-cache` - Bypass cache and make live API calls (falls back to cache on 429 rate limit)
- `--clear-cache` - Clear cached API responses for today's date only
- `--clear-all-cache` - Clear all cached API responses (all dates)
- `--list-cache` - List all cached API responses
- `--inspect-cache <hash|date>` - Inspect cached responses by hash or date (YYYY-MM-DD format)

Examples

```
# Use custom config file
./enphase-monitor --config /path/to/my-config.yaml
```

```
# Query last week's data
./enphase-monitor --once --date 2026-01-12
```

```
# Continuous monitoring with default settings
./enphase-monitor
```

Output Format

The application displays:

```
=====
ENPHASE MULTI-SYSTEM MONITOR
=====
Query Range:    Tue Jan 20, 2026 12:00 AM
                to
```

Tue Jan 20, 2026 11:59 PM

Last Updated: Sat Jan 24, 2026 09:52:18 PM (cached)

COMBINED ENERGY REPORT (kWh)

Produced:	33.4 kWh
Consumed:	48.6 kWh
Net Energy Flow:	19.1 kWh (import)

INDIVIDUAL SYSTEMS REPORT

[1] Right Subpanel (5525881)	
Imported from the Grid:	23.1 kWh
Exported to the Grid:	3.8 kWh
Captured from the Sun:	14.6 kWh
Net Energy Flow:	19.3 kWh (import)
Charged to Battery:	8.5 kWh
Discharged from Battery:	6.8 kWh
Battery Charge Percentage:	63%
Total Consumed:	32.1 kWh
[2] Left Subpanel (5392556)	
Imported from the Grid:	7.5 kWh
Exported to the Grid:	7.6 kWh
Captured from the Sun:	18.9 kWh
Net Energy Flow:	0.2 kWh (export)
Charged to Battery:	8.1 kWh
Discharged from Battery:	5.4 kWh
Battery Charge Percentage:	74%
Total Consumed:	16.4 kWh

Note: Colors are customizable via `config.yaml` (see Color Customization section below).

Metrics Explained

Combined Energy Report

- **Produced:** Total solar generation from all systems (kWh)
- **Consumed:** Total household consumption from all systems (kWh)
- **Net Energy Flow:** Net energy imported from or exported to the grid
 - Positive value with “(import)” suffix: More energy imported than exported
 - Negative value with “(export)” suffix: More energy exported than imported
 - Calculation: Grid Imported - Grid Exported

Individual System Metrics

- **Imported from the Grid:** Energy purchased from utility for this system (kWh)
- **Exported to the Grid:** Energy sold back to utility from this system (kWh)
- **Captured from the Sun:** Solar generation for this system (kWh)
- **Net Energy Flow:** Net import/export for this system (kWh) with (import) or (export) suffix
- **Charged to Battery:** Energy stored in batteries for this system (kWh)
- **Discharged from Battery:** Energy used from batteries for this system (kWh)
- **Battery Charge Percentage:** Current state of charge (SOC) of the battery system, displayed as a percentage (0-100%). This metric is shown per-system only and is not aggregated across multiple systems.
- **Total Consumed:** Total consumption for this system (kWh)

Color Customization

You can customize the colors used in the terminal output by adding a `colors` section to your `config.yaml`:

```
colors:
  production: "#f0b57c"      # Solar Production
  discharge: "#7acf38"       # Battery Discharge
  import: "#f63cb1"          # Grid Import
  export: "#06b6de"          # Grid Export
  net_import: "#f63cb1"      # Net Energy Flow (Import)
  net_export: "#06b6de"      # Net Energy Flow (Export)
  headers: "#f37320"         # Report Headers
  charge: "#7acf38"          # Battery Charge
  total_consumed: "#f37320"  # Total Consumed
  secondary_text: "#808080"  # Secondary Text
  primary_text: "#ffffff"    # Primary Text
  error: "#ff0000"          # Error Text
```

Color Format Options: - **Hex codes** (e.g., `#FF5733`): Automatically converted to ANSI color codes - **ANSI escape codes** (e.g., `\033[38;5;208m`): Used directly as-is

Note: Reset and Bold are constants defined in `constants.go` and cannot be customized.

Troubleshooting

“API configuration required”

- Make sure you have copied `config.yaml.example` to `config.yaml`
- Verify you have filled in `api.key`, `api.client_id`, and `api.client_secret`
- For developer plan, complete OAuth setup with `--setup-oauth` to get `refresh_token`

“API request failed with status 401”

- Your refresh token may have expired or been revoked
- Re-run OAuth setup: `./enphase-monitor --setup-oauth`
- Verify your API credentials are correct

“API request failed with status 404”

- Verify your System IDs are correct (check Enlighten URLs for each system)
- Check that you have access to all systems in Enlighten
- Ensure the system IDs are strings (quoted in YAML)

“rate limit exceeded (429)”

- The API has a rate limit of 10 calls per minute
- The program will display how many seconds to wait before retrying
- Consider increasing `refresh_interval` in your config
- Use `--test` mode to validate against cached data without making API calls

“API request failed with status 422”

- This usually means the requested date is in the future
- The program automatically caps dates to the current time
- Try querying a past date: `--date 2026-01-15`

No telemetry data returned

- Some data may not be available for very recent time periods (try querying yesterday’s data)
- Ensure your systems are actively reporting to Enlighten
- Check cache with `--list-cache` to see what data is available

Caching and Rate Limits

Rate Limits

The Enphase Enlighten Cloud API v4 enforces strict rate limits: - **10 requests per minute** per API key - **1000 requests per month** total (free developer plan)

Refresh Interval Recommendation:

The `refresh_interval` setting controls how often the application queries the API in continuous mode. To respect rate limits:

- **Recommended:** `refresh_interval: 3600` (1 hour)
- **Why:** Each system requires multiple API calls. With 2 systems, you might make 8-10 requests per query cycle. At 3600 seconds, that is ~20 requests per hour, well within limits.
- **Not Recommended:** Values below 60 seconds (e.g., `refresh_interval: 5`) can quickly exceed the 10 requests/minute limit, especially with multiple systems.
- **Calculation:** If you have N systems and each requires M API calls, you make N×M requests per cycle. At `refresh_interval: 5`, that is N×M×12 requests per minute, which can easily exceed 10/minute.

Caching Strategy

To respect these limits, the application implements intelligent caching:

- **Automatic Disk Caching:** All API responses are cached in `cache/` directory

- **Cache-First for Past Dates:** When querying historical dates, cached data is used if available (no API call)
- **Default Refresh Interval:** 1 hour (3600 seconds) - queries each system once per hour
- **429 Error Handling:** If rate limited, the program displays wait time and exits gracefully

Cache File Format and Naming

Cache files are stored in the `cache/` directory with the following structure:

File Naming Scheme: - Each cached API response is stored as a JSON file - Filenames are SHA-256 hashes of the normalized API URL (including query parameters) - Format: `{hash}.json` where `{hash}` is a 64-character hexadecimal string - Example: `a1b2c3d4e5f6...{64 chars}.json`

Cache File Structure: Each cache file contains: - **status_code:** HTTP status code from the API response - **headers:** HTTP response headers (as key-value pairs) - **body:** Raw API response body (JSON bytes) - **cached_at:** Timestamp when the response was cached (ISO 8601 format) - **queried_date:** The date that was queried (YYYY-MM-DD format), if applicable

Cache Key Generation: - Cache keys are generated from normalized API URLs - URLs are normalized to ensure consistent caching (timestamps converted to date strings) - The same API query will always produce the same cache key, regardless of when it's made

Cache Lookup: - When making an API request, the application first checks for a cached response - If a cache file exists and is valid, the cached response is used (no API call) - For past dates, cached data is always preferred to avoid unnecessary API calls

Cache Management Commands

```
# View all cached responses
./enphase-monitor --list-cache

# Inspect a specific cached response (by hash or date)
./enphase-monitor --inspect-cache 2026-01-15

# Clear today's cache only (preserves historical data)
./enphase-monitor --clear-cache

# Clear all cached data
./enphase-monitor --clear-all-cache

# Disable caching (always make live API calls)
./enphase-monitor --no-cache

# Test mode (use cache only, no live API calls)
./enphase-monitor --test --date 2026-01-15
```

Best Practices

1. **Use Default Refresh Interval:** Use 1 hour to stay well within rate limits
2. **Leverage Caching:** Query past dates frequently - they use cached data (no API calls)
3. **Test Mode:** Use `--test` mode for validation against cached data without hitting API

4. **Monitor Cache:** Use `--list-cache` to see what data is available before querying

Documentation

This project includes comprehensive documentation for different learning paths:

For Getting Started

- **QUICKSTART.md** - Get up and running in 5 minutes
- **OAuth_SETUP.md** - Complete OAuth 2.0 setup guide with detailed explanations

For Understanding the Codebase

- **ARCHITECTURE.md** - System architecture, execution flow, and design patterns
- **GO_BEST_PRACTICES.md** - Go concepts and patterns used in this codebase
- **GO_CONCEPTS.md** - Quick reference for Go concepts used in the code, including channels and signals

Recommended Learning Path

1. **New Users:** Start with QUICKSTART.md
2. **OAuth Setup:** Follow OAuth_SETUP.md for authentication
3. **Understanding Go:** Read GO_BEST_PRACTICES.md for Go concepts
4. **System Design:** Study ARCHITECTURE.md to understand the codebase structure
5. **Advanced Topics:** Explore GO_CONCEPTS.md for channels, signals, and concurrency patterns

Project Structure

```
enphase-monitor/  
  main.go                # Application entry point (orchestration only)  
  internal/  
    aggregator/  
      types.go           # Internal packages  
      aggregator.go      # Multi-system data aggregation  
      aggregator_test.go # Metric data structures  
      aggregator_bench_test.go # Aggregation logic with dependency injection  
                                # Aggregator tests with mock clients  
    api/  
      client.go          # Benchmark tests  
      types.go           # HTTP client for Cloud API v4  
      interface.go       # Enlighten Cloud API client  
      client_test.go     # API request/response types  
      client_functional_test.go # CloudClient interface for testability  
                                # API client unit tests  
                                # Functional tests with mock HTTP servers  
  app/  
    setup.go             # Application execution logic  
    setup_test.go        # App initialization & configuration  
    runner.go            # Setup tests  
    runner_test.go       # Execution modes (once/continuous)  
                                # Runner tests  
  cache/  
                                # Disk-based response caching
```

cache.go	# Cache implementation
cache_test.go	# Cache state management tests
cache_functions_test.go	# Cache functionality tests
cli.go	# Cache inspection utilities
cli_test.go	# CLI utilities tests
cli/	# Command-line interface
flags.go	# CLI flag parsing
flags_test.go	# Flag parsing tests
cache_commands.go	# Cache management commands
cache_commands_test.go	# Cache commands tests
config/	# Configuration types
config.go	# YAML loading & validation (uses type aliases)
config_test.go	# Configuration tests
constants/	# Centralized constants
constants.go	# Application-wide constants
constants_test.go	# Constants tests
display/	# Terminal output formatting
display.go	# Display with io.Writer injection
display_test.go	# Display output tests
oauth/	# OAuth 2.0 authentication
oauth.go	# Token management & refresh
setup.go	# Interactive OAuth wizard
oauth_test.go	# Basic unit tests
oauth_functional_test.go	# Integration tests with mock servers
oauth_edge_cases_test.go	# Edge case tests
parser/	# JSON telemetry parsing
parser.go	# Response parsing utilities
parser_test.go	# Parser tests
parser_bench_test.go	# Benchmark tests
timezone/	# Timezone handling
timezone.go	# Timezone utilities
timezone_test.go	# Timezone tests
types/	# Shared type definitions
types.go	# SystemConfig, APIConfig (breaks circular deps)
urlbuilder/	# API URL construction
urlbuilder.go	# URL building helpers
urlbuilder_test.go	# URL builder tests
validation/	# Test mode validation
validation.go	# Metrics validation logic (uses io.Writer for testab
validation_test.go	# Unit tests (tolerance calculations, edge cases)
validation_integration_test.go	# Integration tests (real expected values)
docs/	# Project documentation
ARCHITECTURE.md	# Architecture documentation
GO_BEST_PRACTICES.md	# Go best practices guide
GO_CONCEPTS.md	# Go concepts reference (channels, signals, and more)
OAUTH_SETUP.md	# OAuth setup documentation (detailed)
TESTING.md	# Testing patterns and guidelines
go-documentation-review.md	# Go documentation style review

go-linting-review.md	# Linting configuration review
go-style-core-review.md	# Go style guide compliance review
test-data/	# Test validation data
expected_values_*.json	# Expected values for validation
config.yaml.example	# Example configuration with all options
config.yaml	# Your actual configuration (create from example)
cache/	# Cached API responses (created at runtime)
go.mod	# Go module definition
go.sum	# Go module checksums
scripts/	# Utility scripts
generate-pdfs.sh	# Script to generate PDFs from markdown files
run-tests.sh	# Test runner script
Makefile	# Build automation
README.md	# This file
QUICKSTART.md	# Quick start guide
pdfs/	# Generated PDF documentation (created by generate-pdf)

Testing

The project includes a comprehensive test suite with **70.4% code coverage** across all packages. The test suite validates both functionality and metrics against expected values, enabling rapid iteration without hitting API rate limits.

Test Coverage by Package

Package	Coverage	Status
constants	100.0%	
display	100.0%	
urlbuilder	100.0%	
validation	96.6%	
timezone	93.3%	
config	82.4%	
parser	80.8%	
aggregator	80.0%	
app	76.8%	
api	74.4%	
cache	66.9%	
cli	47.6%	
oauth	46.1%	

Total: 70.4% coverage (exceeds typical Go project standards of 50-60%)

Running Tests

```
# Run all tests
go test ./...
```

```

# Run with coverage report
go test ./... -coverprofile=coverage.out
go tool cover -html=coverage.out

# Run specific package tests
go test ./internal/parser -v

# Run benchmarks (see Performance section below)
go test -bench=. ./internal/parser
go test -bench=. ./internal/aggregator

```

Lint and CI

Run the linter (errcheck, goimports, revive, govet, staticcheck) before committing:

```
make lint
```

CI (GitHub Actions) runs on push and pull requests: it builds the project, runs `go vet` and tests, and runs `golangci-lint`. See `.github/workflows/ci.yml`.

Test Mode (Cache Only)

Run in test mode using only cached responses (no live API calls):

```
./enphase-monitor --once --test --date 2026-01-20
```

This will: 1. Use only cached API responses (no live calls) 2. Validate results against expected values 3. Show a detailed comparison report

Early Validation The `--test` flag includes early validation to prevent confusing errors:

Missing cache:

```
ERROR: --test flag requires cached data, but no cache exists for 2026-01-30.
```

To populate the cache, run:

```
./enphase-monitor --once
```

Then retry with `--test`.

Missing expected values file:

```
Validation failed: no expected values file found for 2026-01-01.
```

To run validation, create the file:

```
test-data/expected_values_2026-01-01.json
```

Example format:

```
{ "date": "2026-01-01", "systems": [...] }
```

To skip validation and just use cache-only mode, omit the `--test` flag:

```
./enphase-monitor --once --date 2026-01-01
```

Setting Up Test Data

IMPORTANT: The cache must contain data for the specific test date. Old cache files from different dates will cause incorrect results.

1. **Run all tests** (recommended):

```
./scripts/run-tests.sh
```

This script will:

- Check which test dates have cached responses
- Generate missing cache by making live API calls (waiting 60 seconds between calls to respect rate limits)
- Run validation tests for all dates (2026-01-14 through 2026-01-20)
- Display a summary of all test results

2. **Or manually generate cache for a specific date:**

```
./enphase-monitor --once --date 2026-01-20
```

This will make live API calls and cache the responses in `cache/`

3. **Update `expected_values_YYYY-MM-DD.json`** with the correct values from the Enphase app
4. **Now you can iterate rapidly using test mode** (uses cache only, no API calls):

```
./enphase-monitor --once --test --date 2026-01-20
```

Note: Test mode uses the cache from `cache/`. The `scripts/run-tests.sh` script ensures all test dates have cached responses available.

Expected Values Format

The expected values JSON file should have this structure:

```
{
  "date": "2026-01-20",
  "systems": [
    {
      "id": "5525881",
      "name": "Right Subpanel",
      "expected": {
        "grid_import": 23.4,
        "grid_export": 3.9,
        "production": 14.6,
        "battery_discharged": 6.8,
        "battery_charged": 8.6,
        "net_imported": 19.6,
        "consumption": 32.3
      }
    }
  ]
}
```


Validation Tolerance

The validation function compares actual vs expected values with a tolerance of **$\pm 10\%$ of the expected value**, with a minimum tolerance of **0.1 kWh** for small values. This means: - For large values (e.g., 20 kWh): tolerance is ± 2.0 kWh (10%) - For small values (e.g., 0.5 kWh): tolerance is ± 0.1 kWh (minimum)

Metrics that differ by more than the calculated tolerance will be marked as failed (). The validation report shows: - Expected value - Actual value - Difference - Percentage difference - Pass/fail status

Performance

The codebase includes comprehensive benchmarks for performance-critical paths. Run benchmarks with:

```
# Parser benchmarks (JSON parsing, interval summing)
go test -bench=. -benchmem ./internal/parser

# Aggregator benchmarks (multi-system aggregation)
go test -bench=. -benchmem ./internal/aggregator

# All benchmarks with CPU profiling
go test -bench=. -benchmem -cpuprofile=cpu.prof ./internal/...
```

Key Performance Characteristics: - **API Response Caching:** Reduces redundant API calls and respects rate limits - **Efficient JSON Parsing:** Optimized for 96 intervals/day (15-min intervals) - **Multi-System Aggregation:** Scales linearly with number of systems - **Zero External Dependencies:** Standard library only for performance and security

Code Quality

This project follows Go best practices and coding standards:

- **Test Coverage:** 70.4% overall, 100% for critical packages (constants, display, urlbuilder)
- **Test Suite:** 24 test files across 13 packages with comprehensive unit, integration, and edge case tests
- **Go Modules:** Proper dependency management with go.mod/go.sum
- **Error Handling:** Comprehensive error wrapping with context
- **Documentation:** Extensive inline comments and dedicated guides
- **Type Safety:** Strict type checking with no unsafe operations
- **Linting:** Passes golangci-lint with recommended settings
- **Performance:** Benchmarks included for hot paths

Code Metrics: - Total Lines: ~4,000 (excluding tests) - Test Lines: ~3,500+ (comprehensive test suite) - Packages: 13 internal packages - Test Files: 24 (unit, integration, functional, edge case, and benchmark tests) - External Dependencies: 1 (gopkg.in/yaml.v3)

License

This is a personal utility project. Use and modify as needed for your own Enphase monitoring needs.